
Practical Machine Learning: Pipeline Optimization

Russell Todd¹

¹University of Wyoming

1 Introduction

In this exercise, I will be looking at a more sophisticated version of automated machine learning and optimizing the entire pipeline which includes preprocessing, the type of machine learning model, and the model hyperparameters for my task. I will familiarize myself with the process and report on how I progressed through the exercise. I will strive to explain well enough that reproducing my work would be doable even if I did not share my exact code. The dataset I will be looking at is the White Wine Quality dataset so I will be using the features of wines to predict their quality score.

2 Dataset Description

The famous White Wine Quality dataset is comprised of 12 features. The first 11 are measurements of various physical qualities of each wine (fixed acidity, citric acid, residual sugar, pH, alcohol, etc.) and the 12th is the output variable which is a score (0-10) denoting the quality of each wine. There are 4,897 entries (wines) in the dataset and no entries are missing data for any feature. All features are positive numeric values.

After exploring the White Wine Quality dataset, it can be seen that it is quite imbalanced with relatively few entries for low quality and high quality wines. Therefore I will try using oversampling techniques to correct this.

It also has two variables, density and residual sugar, that are highly correlated. I decided to drop density before training models as it is less correlated to the score than residual sugar.

Additionally, most of the features have significant outliers present in their data as can be seen from a selection of plots in figure 3.

3 Experimental Setup

The optimization pipeline encompasses preprocessing, model selection, and hyperparameter optimization with each step explained below. The first step was to get the data properly loaded into a pandas package dataframe and ensure that the datatypes of each feature are appropriate.

Afterwards, several summary statistics panes are produced using built-in functions of the pandas dataframe class like `info()`, `dtypes`, and `describe()`. I then use these to double check that the data loading process was done correctly. If not, I then corrected them manually.

The selected regression algorithms were LogisticRegression, DecisionTreeClassifier, RandomForestClassifier, and a Radial Basis Function Support Vector Machine (RBF SVM). All algorithms were first run with parameters as close to default as possible, with the exception of LogisticRegression which required me to raise the max iterations parameter to run it properly.

To assess the models performance I used the accuracy score from the sklearn metrics package on a fivefold cross-validation run of each model to try to avoid overfitting. This basic output was then stored for later comparison to the optimized models after hyperparameter optimization.

3.1 Model Selection with Hyperparameters Step

Using the pipeline to streamline the process, I first split the data into a an 80:20 training test split before entering the preprocessing step. The preprocessing is then handled by the pipeline which

uses the "preprocessor__num__with__mean" hyperparameter with the binary options of True or False to control the Standard Scaler from the sklearn preprocessing package for numeric features.

For the pipeline, the next step was to define the classifier hyperparameter search space. I chose the range for each of the hyperparameters from reading the documentation on each model's parameters and selected what I felt were reasonable numbers (basically just multiples of 10).

The hyperparameters that were tuned for each model are as follows: preprocessor__num__with__mean with a binary range of True or False, classifier__C with a Real range of 0.1 to 10 for LRC. classifier__max_depth with an Integer range of 10 to 100 for DT. classifier__n_estimators with an Integer range of 100 to 1000 and classifier__max_depth with an Integer range of 10 to 100 for RFC. classifier__C with a Real range of 0.1 to 10 and classifier__gamma with a Real range of 0.1 to 5.

The driving force behind exploring the search space was a bayesian search evaluating performance based on accuracy scores from fivefold cross-validation. The bayes search was implemented using the sklearn optimization package (skopt). To prevent overfitting, nested resampling was performed as part of the bayes search using a 5-fold cross-validation strategy. The bayesian optimization was run for 100 iterations. The scores were stored in much the same way as the basic output scores were so that I see the performance difference after the optimization. This data included the training and testing data so that potential overfitting could be seen.

4 Results

From the graphs below it can be seen that the RandomForestClassifier noticeably outperformed the other competing algorithms which somewhat surprised me as the RBFSVM is supposed to be very sophisticated. However, the RBFSVM did show the greatest improvement in performance with optimization so perhaps with even more time it could surpass the RFC.




Figure 1: Basic Classification Model Scores




Figure 2: Optimized Classification Model Scores




Figure 3: Effect of Optimization on Classification Models

The exact scores are shown below:

Basic Scores

LRC Train Accuracy: 0.5400669846482654
LRC Test Accuracy: 0.5183673469387755

DT Train Accuracy: 0.5862606536867621
DT Test Accuracy: 0.6081632653061224

RFC Train Accuracy: 0.6618145150780619
RFC Test Accuracy: 0.7081632653061225

RBFSVM Train Accuracy: 0.5758089634321162
RBFSVM Test Accuracy: 0.5571428571428572

Optimized Scores

LRC Train Accuracy: 0.5431291865405166
LRC Test Accuracy: 0.5255102040816326
LRC Optimal Parameters:
OrderedDict([('classifier__C', 0.1136442723231016), ('preprocessor__num__with_mean', False)])

DT Train Accuracy: 0.5875361638908437
DT Test Accuracy: 0.6081632653061224
DT Optimal Parameters:
OrderedDict([('classifier__max_depth', 83), ('preprocessor__num__with_mean', True)])

RFC Train Accuracy: 0.6658958219303047
RFC Test Accuracy: 0.7030612244897959
RFC Optimal Parameters:
OrderedDict([('classifier__max_depth', 76), ('classifier__n_estimators', 945), ('preprocessor__num__w

RBFSVM Train Accuracy: 0.6342507884379806
RBFSVM Test Accuracy: 0.6642857142857143
RBFSVM Optimal Parameters:
OrderedDict([('classifier__C', 3.6083793798349446), ('classifier__gamma', 0.9213778732677528), ('prep

All optimized models had scores greater than or equal to their basic counterparts with the exception of the RandomForestClassifier.